

Solving the Maximum Subsequence Problem with a Hardware Agents-based System

Octavian Creț¹, Zsolt Mathe¹, Cristina Grama¹, Lucia Văcariu¹, Flaviu Roman¹, Adrian Dărabant²
Computer Science Department

¹ Technical University of Cluj-Napoca; ² “Babeș-Bolyai” University of Cluj-Napoca
26, G. Barițiu Street, Cluj-Napoca; 1, M. Kogălniceanu Street, Cluj-Napoca
ROMANIA

{Octavian.Cret , Zsolt.Mathe, Cristina.Grama, Lucia.Vacariu, Flaviu.Roman}@cs.utcluj.ro

Abstract – The maximum subsequence problem is widely encountered in various digital processing systems. Given a stream of both positive and negative integers, it consists of determining the subsequence of maximal sum inside the input stream. In its two-dimensional version, the input is an array of both positive and negative integers, and the problem consists of determining the sub-array of maximal sum inside the input array. These problems are solved by Kadane’s algorithm, which has already been proved to be optimal. However, the hardware implementation presented in this paper is based on the hardware agents’ paradigm and offers a significantly improved performance (in terms of speed) over the classical software implementations.

Key-words: Kadane algorithm, maximal subsequence problem, hardware agents, VHDL, FPGA

1 Introduction

The problem of maximum sum was first introduced by Bentley [1, 2] and has one-dimensional and two-dimensional versions. The 1D version is referred to as the maximum subsequence problem, while the other is called the maximum sub-array problem.

The maximum subsequence problem is very widely encountered in various digital processing systems. It consists of determining a contiguous portion of an unsorted 1D array that contains both positive and negative integers in which the sum of the array elements is maximal, over the entire stream.

For instance, in the following 13-element 1D array:

$a[13] = \{7, -9, 15, 20, -37, 23, 4, 5, 19, -28, 17, -2, 1\}$

If we consider the first element to be $a[0]$, the maximal sub-array is delimited by the element of index 5 ($a[5] = 23$) and the element of index 8 ($a[8] = 19$).

This problem does not appear to be a very difficult one to implement in software. It is implemented for various DSP applications [6]; Bentley [1] introduced Kadane’s algorithm that finds the maximum subsequence of a 1D array in linear time $O(n)$.

In the following, let s be the notation for the sum of a tentative maximum subsequence $a[i..j]$. The algorithm accumulates a partial sum in t and, when t becomes better than s , it replaces s with t and updates the position (i, j) . If t becomes negative, we reset the accumulation.

This operation is due to its dynamic programming nature; a subsequence of negative sum can only constitute a suboptimal maximum sum, thus it is discarded. The algorithm is presented in Fig. 1.

```
1.  input :  $a[1..n]$ 
2.  output :  $s, (x1, x2)$ 
3.   $(x1, x2) \leftarrow (0, 0); s \leftarrow -\infty; t \leftarrow 0;$ 
4.   $i \leftarrow 1;$ 

5.  for  $j \leftarrow 1$  to  $n$  do
6.     $t \leftarrow t + a[j];$ 
7.    if  $t > s$  then
8.       $(x1, x2) \leftarrow (i, j); s \leftarrow t$ 
9.    end

Reset the accumulation;
10. if  $t < 0$  then
11.    $t \leftarrow 0; i \leftarrow j + 1$ 
12. end
13. end
```

Fig. 1. Kadane’s algorithm

One of the major tendencies in digital computing systems is to relocate the computation-intensive parts of applications into hardware.

Processors have a general, fixed architecture that allows the implementing of tasks by *temporally* composing atomic operations provided e. g. by the ALU or the floating-point unit. In contrast, ASICs implement tasks by *spatially* composing operations provided by dedicated functional units. Reconfigurable computing combines both approaches. Reconfigurable systems are built from programmable logic, which allows the implementing of tasks both in a spatial manner similar to ASICs and in a temporal manner comparable to processors.

The main goal of our research group is to develop hardware implementations for the most widely used

classical software algorithms. Thus, it will be possible to gradually implement larger systems based on specialized cooperative hardware agents, where each agent is responsible for implementing a given algorithm.

We use an agent-oriented approach, considering an agent to be an intelligent and, to a certain degree, autonomous, entity which has a certain skill (in our case, the skill is the ability to use the hardware implementation of an algorithm to solve an instance of that algorithm). An agent can make use of its skill to solve a particular problem. However, in practice, a single agent cannot attain the major goal of the system by itself – the agents need to cooperate in order to achieve the final goal. Furthermore, we believe the hardware-implemented agents will have an edge over classical software agents, especially what the speed of the system is concerned.

This paper shows the results for the implementation of Kadane's algorithm in software and in hardware, focusing on the hardware implementations which yield much better results. The design technique is detailed, and the necessary adaptations are explained. Finally, the architecture is loaded into a Xilinx FPGA device, for simulation and testing purposes.

The paper is structured as follows. After this Introduction, Section 2 presents the results of the software implementation for the 1D case of the algorithm. The design of the 1D hardware architecture is detailed in Section 3. A comparison between the hardware and software implementations for the 1D case is presented in Section 4. Section 5 treats the 2D case from both the software and the hardware point of view. Results and final Conclusions are presented in Section 6.

2 The Software Implementation in the 1D Case

The software implementation for the 1D case is straightforward and does not deserve to be detailed here. The program was written in C++ and tested on a few PC systems with Pentium IV processors running Mandrake Linux 10.1, with Linux Kernel 2.6.8. The obtained results are presented in Table 1.

It must be noticed that these results are due to the *temporal* organization of the computing task in the microprocessor-based architecture. They are also limited by the numerous memory accesses required by this algorithm.

As it can be seen in Table 1, the run time for a 60,000 numbers sequence is of approximately 90-100 ms. In the third column of this table we can check the running times of the optimal software implementation in seconds. The first column contains the number of times the module was run on the input stream, whose length is

given in the second column. It is advised [7] to do more tests for the same input length to assure stabilized results.

Table 1. Results of the software implementation

Number of Tests	Input Stream Length	Run time in C (seconds)
1	100000	0.231
10	10000	0.16
10	20000	0.3
10	30000	1.012
100	10000	1.722
100	20000	3.385
100	30000	5.208
1000	10000	17.805
1000	20000	37.013
1000	30000	45.956
1000	60000	93.789

The much better results achieved by the hardware implementation are due to the spatial organization of the computing task in an FPGA device, which has been used as physical support for the implementation.

3 The Hardware Implementation in the 1D Case

First the algorithm must be adapted for a hardware implementation. We consider the case of the analyzed stream received serially, synchronously, on our system's data input. The outputs generated by the design are the two pointers to the beginning and the end of the maximal sequence, x_1 and x_2 , and the maximal sum obtained, s . Fig. 2 presents the block diagram of the design.

As Naji & Wells have shown in [4], hardware multi-agent systems allow the concepts of multi agent technology to be extended to support reconfigurable systems, in which the functionality of both the associated hardware and software can be altered dynamically or statically with time. The utilization of such a paradigm has the potential to greatly increase the flexibility, efficiency, expandability, and maintainability of such systems and to provide an attractive alternative to the current set of disjoint approaches that are currently applied to this problem domain [4, 5].

This paper presents an example of how an agents-based architecture can be created and applied to the reconfigurable hardware domain.

We'll consider the system at a certain moment in time to be an *agent*. This agent's beliefs are the following: temporary maximal sum s , partial sum t , current data from vector $a[j]$ (DATAIN(a)), and temporary indices i and j .




Fig. 2. Block diagram of the hardware system for Kadane's algorithm

Its goals are: the left index of the maximal sum sub-array x_1 , the right index of the maximal sum sub-array x_2 , and the temporary (updated) sums s and t .

The agent's plan is the actual algorithm (Fig. 3):

```

Inputs = (s, t, a[j], i, j)
ag_mem[0] ← Input(0)
ag_mem[1] ← Input(1)
ag_mem[2] ← Input(2)
ag_mem[3] ← Input(3)
ag_mem[4] ← Input(4)
ag_mem[1] ← ag_mem[1] + ag_mem[2]
if (ag_mem[1] > ag_mem[0]) then
  (ag_mem[5], ag_mem[6]) ← (ag_mem[3], ag_mem[4])
  ag_mem[0] ← ag_mem[1]
end
if (ag_mem[1] < 0) then
  ag_mem[1] ← 0
  ag_mem[3] ← ag_mem[1] + 1
end
Outputs = (ag_mem[0], ag_mem[1], ag_mem[5],
ag_mem[6]) // (s, t, x1, x2)

```

Fig. 3. The agents-based algorithm implementation

When passing from software to hardware, there are a few additional elements to take into consideration.

3.1 Input Data, Accumulators and Adder

As indicated in Section 1, the numbers from the input stream are signed integers. It is thus necessary to design the whole architecture in two's complement arithmetic.

A common sense estimation (which also covers a wide range of practical situations) is to consider the input numbers as bytes. Therefore, we need an 8-bit representation of these numbers, and one extra bit for the sign. This is why, in Fig. 2, the DATAIN port is sized as a 9-bit input.

In the 1D case, our purpose was to design a system that works properly for 65,536 numbers in the input stream. Since each number is expressed on 8 bits, a simple estimation shows that, for the worst case situation, we need an *Accumulator* on 25 bits to compute the partial sums (t , in the algorithm presented in Fig. 1).

Also, the temporary maximal sum, s , must also be sized as a 24-bit data register, because s will never be negative, as one can easily deduce from the algorithm, so the sign bit is no longer necessary.

Of course, an *Adder* resource is necessary. This resource does not appear explicitly in Fig. 1, because

Kadane's algorithm was targeted for a software implementation. The *Adder* too must be sized on 24 bits.

3.2 The j Iterator, Internal Pointers and Variables

As can be seen in Fig. 1, the j variable iterates from 1 to n (n being the size of the input stream). In the hardware implementation, this corresponds to a *Counter*. Our goal was to design a system that deals with streams of at most 65,536 numbers, so the *Counter* was sized on 16 bits.

The algorithm described in Fig. 1 also contains three more variables:

- x_1 , the pointer to the beginning of the maximal subsequence found in the input stream
- x_2 , the pointer to the end of the maximal subsequence found in the input stream
- i , an intermediate variable that points to the candidate maximal subsequence.

These variables are stored into data registers. All of them must be sized on 16 bits.

3.3 The Increment (“+ 2”) Block

When moving the algorithm from software to hardware, some additional resources are necessary, which do not appear explicitly in the pseudo code: the *Adder* and the *Counter* are examples of such “hidden” resources. Another hidden resource is the Increment (“+ 2”) Block.

The Increment (“+ 2”) Block is necessary for implementing the (apparently) simple operation

$$i \leftarrow j + 1$$

in the last line (line 11) of the algorithm.

The goal of the hardware implementation is to obtain a speed increase. The speed gain becomes significant for large sizes of the problem (in this case, for a greater number of elements in the input stream). For this purpose, our intention was to create a pipelined design, where in each clock cycle a new number from the input stream is processed. The whole design must of course be synchronous.

It is generally a good idea to buffer the input data, so, as can be seen in Fig. 4, a new partial sum ($t \leftarrow t + a[j]$) will be obtained one clock cycle later than the arrival of a new input number.

This organization also has an advantage. At the end of the algorithm (lines 10-12 on Fig. 1), we have to compare t to 0. Because of this structure (see Fig. 4), we know in advance if the *Accumulator* (that stores t) will be loaded with a negative number. In software this operation would take place exactly as described by the algorithm: first, load t with the new value ($t \leftarrow t + a[j]$), then test if it is negative. In hardware, since t is a data register loaded with the result produced by the *Adder*, if the first bit is ‘1’ (meaning that the result of the addition is negative), on the next clock cycle we can directly and synchronously reset the *Accumulator*.




Fig. 4. The hardware architecture in the 1D case

Because of this structure, the algorithm must be adapted as follows. At the end of the current iteration, the i variable must be updated to the value $j + 1$. In our case, we must perform the operation:

$$i \leftarrow j + 2$$

because t “will be negative” only later, after one clock cycle. But after that clock cycle, j will be incremented, so i must be updated to the value $j + 2$.

3.4 The Comparator and the Sign Extension

The Comparator is a combinational resource needed to determine if $t > s$ (line 7 in the algorithm shown in Fig. 1). It must be sized on 24 bits, because s is always positive and is compared to t only when t is positive too, so for t the sign bit is no longer necessary.

Technically speaking, the sign extension isn't a resource, but it is implemented by the interconnection network. It is necessary to perform a sign extension on the input numbers, since they come as two's complement 9-bit words, while the Adder is sized on 25 bits.

To preserve the readability of the schematic, in Fig. 4 the *Clock* and *Preset* signals have not been represented, for all the sequential components (the 9-bit Input Buffer, the 25-bit Accumulator, the 24-bit Data Register, the 16-bit Data Registers, and the Counter).

3.5 Initialization

The Initialization of the designed system is done as follows. All the sequential components are initialized with 0, except for the 16-bit Counter that stores j , which will be initialized with $2^{16}-2$, i.e. “111111111111110”. The reason is given by the particular features of this architecture, as explained above.

4 Results of the 1D Implementation

The implemented algorithm is the most simple and optimal solution for the maximum subsequence problem and it can be used in the solution of bi-dimensional problem (the solution for the problem also known as the maximum sub-array problem relies on partial subsequence calculus).

A small modification of Kadane's algorithm is used by [3] for solving the K maximum subsequence, using the concept of prefix sums. While this second algorithm works just as well as Kadane's algorithm both in hardware and software implementation, in the implementation for the 2D solution the hardware implementation of the second algorithm is rather limited to small data inputs (because of the need of special arrays keeping the prefix sums).

A detailed methodology for adapting the 1D Kadane algorithm for a hardware implementation was presented. All choices made for the nature and the size of the used resources are discussed and justified.

The hardware design was specified in parameterizable VHDL code and the FPGA chip that was chosen as physical support of the implementation was a Xilinx Virtex II 1000, speed grade -6 device.

After running Xilinx Synthesis software, the reports indicated a maximal working frequency of 133 MHz (clock period: 7.5 ns). The simulation was done in ModelSIM and the proper functioning of the design was confirmed.

It is obvious that the performance depends on the length of the input stream. A simple computation/calculation indicates that for an input stream of 65,536 numbers, the total execution time is 0.49 ms, which indicates at least two orders of magnitude improvement over the software implementation.

An examination of the results yielded by the software implementation shows that the run time does not increase linearly with the length of the input stream, but in case of the hardware implementation, this dependency is linear.

The FPGA device is used only at 20% of its total capacity, which allows the implementation of the hardware version for much larger streams of input numbers. It is obvious that the performance depends on the length of the input stream.

As a result, the more we increase the length of the input stream, the better the hardware implementation performs when compared with the software implementation.

It is easy to adapt the hardware architecture to larger sizes of the input stream by resizing the resources on a larger number of bits, which is easy to do in VHDL, due to the usage of parameterizable code (based on generics).

5 The Implementation in the 2D Case

The 2D variant of this problem consists of finding the sub-array of maximal sum inside a given $M \times N$ array. For example, for the array in Fig. 5, the solution is highlighted in gray.

2	-4	5	-5	-9	3	-4
-5	-6	-9	2	-5	4	1
-3	7	-7	5	6	-6	0
-3	5	-8	8	-2	2	-7
-4	1	-1	0	-1	-4	1

Fig. 5. Example of the 2D problem solution

The software solution to this problem is classical and straightforward to implement in a program. The results are shown in Table 2.

Table 2. Results of the software implementation

Nr of tests	Input Array Size ($M = N$)	Run time in C (seconds)
10	4	0.0062
1000	4	0.6005
10	16	0.0282
1000	16	2.7902
10	64	0.0156
1000	64	1.5050
10	128	0.0516
1000	128	5.0300
10	170	0.1600
100	170	1.5900
1000	170	15.5500
1	512	1.7030
10	512	16.6300
1	648	3.3910
10	648	33.4400
1	1024	13.187
10	1024	128.870
1	2048	99.218

In the third column of the above table we can check the running times of the optimal software implementation of the Kadane2D in seconds. The first column contains the number of times the module was run on the input stream, whose length is given in the second column. It is advised to do more tests for the same input length[8] to assure stabilised results.

The 2D Kadane algorithm is based on the 1D Kadane algorithm. The 1D algorithm is run on the elements of each row of the array ($row_1, row_2, \dots, row_M$), considered as a 1D stream, then, on the sum of each pair of rows ($row_1 + row_2, row_1 + row_3, \dots, row_{M-1} + row_M$). The solution is given by the maximal sum produced by the 1D Kadane algorithm on these cases. If x_1 and x_2 are the pointers to the beginning and the end of the maximal sub-stream, and R_i and R_j are the two added rows for which the sum is maximal, then the solution is delimited by the rectangle given by the upper-left corner (R_i, x_1) and the lower-right corner (R_j, x_2).

As one can see in Table 2, the execution time increases exponentially, $O((M \times N)^3)$ with the size of the input array. The hardware implementation starts showing its usefulness for large arrays, where the gain of speed becomes significant.

For the hardware implementation, a series of adaptations of the algorithm had to be done, in order to wisely exploit the logic resources of the Virtex FPGA device. The input array is stored in a memory. Depending on the space available in the FPGA device, this memory can be implemented inside or outside the chip (in the Virtex BlockRAMs or in an external RAM). Data are read from this memory, word by word, and the hardware architecture is adapted for a pipeline processing style.

5.1 The Agents-based Approach

The agents will eventually be represented by entities and associated processes, so the inputs Inp (or beliefs) and outputs Outp (or the goals) of the hardware agents will be the ports of these entities. The **AgentName()** pseudo-call simulates "invoking" the agent, or modifying the ports to which its associated process is sensitive.

In section 2, we only had one agent, whose goal was to perform the computations in a single loop of the Kadane 1D algorithm. In order to use the **Kadane1D** agent as a part of the Kadane2D-related computations, we need to add a *for* loop wrapper (we'll actually use a *while* loop), so the **Kadane1D** agent looks like this:

```

Kadane1D
Inputs = (n, a[1], a[2], ..., a[n])
ag_mem[0] ← 0 //x1
ag_mem[1] ← 0 //x2
ag_mem[2] ← -30000 //s
ag_mem[3] ← 0 //t
ag_mem[4] ← 1 //i
ag_mem[5] ← 1 //j
ag_mem[6] ← Inputs(0) // n
while (ag_mem[5] < ag_mem[6]) do
    Kadane1DLoop.Inputs(ag_mem[2], ag_mem[3],
Input[ag_mem[5]], ag_mem[4], ag_mem[5])
    // "invoke" the Kadane1DLoop agent:
    Kadane1DLoop()
    ag_mem[2] ← Kadane1DLoop.Outputs(0)
    ag_mem[3] ← Kadane1DLoop.Outputs(1)
    ag_mem[0] ← Kadane1DLoop.Outputs(2)
    ag_mem[1] ← Kadane1DLoop.Outputs(3)
    ag_mem[5] ← ag_mem[5] + 1 // increment
end while
Outputs = (ag_mem[2], ag_mem[0], ag_mem[1]) //s,
x1, x2

```

Fig. 6. The updated Kadane1D agent

We'll model the **Kadane2D** agent starting from the hardware components it consists of, considering each of them an agent which will be invoked by **Kadane2D**:

- **MAX**: Inputs = (s, x1, x2);
Outputs = (maxs, x1, x2, r1, r2)
- **Kadane1D**: Inputs = (a[]); Outputs = (s, x1, x2)
- **MEM**: Inputs = (ADR); Outputs = (DOUT)
- **AddressGen**: Inputs = (N, M); Outputs = (ADR)

Besides several variables, the **Kadane2D** agent also needs to store an array of data at each step. As can be seen in Fig. 7, this array is an actual row from the input matrix, or a sum of such rows; and taking into account the fact that the input matrix has N lines and M columns, the arrays' dimension will be M .

The **Kadane2D** agent's auxiliary storage space (**aux** in the pseudo code) will then be a matrix like this:

	0	1	2	3	4	5	6	7	8	...	M-1
0	N	M	ADR	S	X1	X2	R1	R2	I		
1											
2											

Fig. 7. The **Kadane2D** agent's auxiliary memory

Row 0 stores variables; row 1 stores the DOUT array, whereas row 2 will store the ROWBUF array. Thus, we have a new plan for the Kadane2D agent:

```

Kadane2D

Inputs=(N, M, a[])
aux[0,0] ← Inputs(0)
aux[0,1] ← Inputs(1)

// "Invoke" the address generator module:

AddressGen.Inputs(aux[0,0], aux[0,1])
AddressGen()
// And get the address:
aux[0,2] ← AddressGen.Outputs(0) // ADR

// "Invoke" the memory module:

MEM.Inputs(aux[0,2])
MEM()
// Get DOUT!
aux[0,8] ← 0
while (aux[0,8] < aux[0,1]) do // i<m
    aux[1,aux[0,8]] ← Mem.Outputs(aux[0,8])
    aux[0,8] ← aux[0,8]+1
end while

// Get ROWBUF!

aux[0,8] ← 0
while (aux[0,8] < aux[0,1]) do // i<m
    // ROWBUF ← DOUT + ROWBUF
    aux[2,aux[0,8]] ← aux[1,aux[0,8]]+aux[2,aux[0,8]]
    aux[0,8] ← aux[0,8]+1
end while

```

// Kadane1D's inputs are the values in the ROWBUF array:

```

aux[0,8] ← 0
while (aux[0,8] < aux[0,1]) do // i<m
    KADANE1D.Inputs(aux[0,8]) ← aux[2,aux[0,8]]
    aux[0,8] ← aux[0,8]+1
end while

// "Invoke" Kadane1D
KADANE1D()
aux[0,3] ← KADANE1D.Outputs(0) //s
aux[0,4] ← KADANE1D.Outputs(1) //x1
aux[0,5] ← KADANE1D.Outputs(2) //x2

MAX.Inputs(aux[0,3], aux[0,4], aux[0,5])

// "Invoke" MAX
MAX()
aux[0,3] ← MAX.Outputs(0) //maxs
aux[0,4] ← MAX.Outputs(1) // x1
aux[0,5] ← MAX.Outputs(2) // x2
aux[0,6] ← MAX.Outputs(3) // r1
aux[0,7] ← MAX.Outputs(4) // r2
Outputs = (aux[0,3], aux[0,4], aux[0,5], aux[0,6], aux[0,7])

```

Fig. 8. The **Kadane2D** agent

5.2 The Address Generator

In order to extract rows from the memory the way mentioned above, a special Address Generator module had to be designed. It is presented din Fig. 9.




Fig. 9. The Address Generator

In order to minimize the space occupied inside the FPGA chip, the Multiplier from Fig. 9 is a *MULT18x18* Virtex2 primitive (*DSP48*, in case a Virtex4 device is used). Also, to implement the memory, Virtex BlockRAMs were used. This way, the speed of the design increases because specialized components are used, and also the number of available Virtex slices makes it possible to implement large designs.

The Address Generator has the very important task of generating the appropriate addresses for the Memory to ensure the proper functioning of the algorithm. It must first extract Rows from 0 to $M-1$, then Rows from 1 to $M-1$, then Rows from 2 to $M-1$... etc., until Rows from

$M-2$ to $M-1$. Combined with the architecture of the RowBuffer and helped by the Command Unit (Fig. 10 and Fig. 11), it will help feeding into the Kadane1D module each Row individually, then the sum of each pair of Rows from the input array. This functionality is ensured by the use of the Multiplier from Fig. 9, and also the two Counters.

5.3 The Row Buffer

The RowBuffer is a special FIFO module that allows to extract rows from the memory and to add them. Thus, a new word (array element) is extracted in each clock cycle, in a pipeline fashion. The size of this component is given by the number of columns in the array multiplied by the width of each word (in this case, a nine-bit representation).




Fig. 10. Structure of the RowBuffer

5.4 The Kadane1D Module

An important part of the design is the Kadane1D module, which operates on individual rows or on the result produced by summing pairs of rows, in the order specified above.

Here, an adaptation of this module is necessary, modifying the version presented at the beginning of this paper (see Fig. 2 and Fig. 4), because by adding the Rows the size of the words increases. So, the input of the Kadane1D module will not be on 9 bits any more, but on $9 + \log_2(M \times N)$ bits. This change can also be observed in Fig. 10.

5.5 The MAX Module and the Command Unit

The output of the Kadane1D module is fed into the MAX module, which will determine the maximal sum, together with some other vital information: the current Row or pair of Rows, the x_1 and x_2 pointers.

The Command Unit controls the functioning of the whole system. It is a relatively simple Finite State Machine (FSM) whose operating mode can be described as follows (its state-diagram is shown in Fig. 11):

First, a Row (let's call it Row_i) is extracted from the Memory into the RowBuffer. Due to the circular register-like organization of the RowBuffer, each element of the current Row is added to its corresponding element in the next Row (called here Row_{i+1}). Simultaneously, the elements of Row_i are fed into the Kadane1D module. This way, two tasks are performed in parallel:

- 1) The maximal sub-stream of Row_i is determined in the Kadane1D module;
- 2) The sum between each element of Row_i and the corresponding element from Row_{i+1} is fed into the RowBuffer, ready for the next phase.




Fig. 11. The state-diagram of the Command Unit

The main role of the Command Unit is to stop (HOLD) the functioning of the main modules and synchronously reset those who need to be re-initialized.

After a new Row (or the sum of two Rows) is loaded into the RowBuffer, the Command Unit will HOLD the Memory, the Address Generator and the RowBuffer until the Kadane1D module finishes determining the maximal subsequence inside the stream Row_i . As can be seen from Sections III and IV, one more clock cycle is necessary. Then, the result produced by the Kadane1D module is transmitted to the MAX module.

If Row_j is Row_{M-1} , this means the Command Unit also has to reset the RowBuffer (until this moment, we obtained Row_i , $Row_i + Row_{i+1}$, ..., $Row_i + Row_{M-1}$, but now we have to obtain Row_{i+1} , $Row_{i+1} + Row_{i+2}$, ..., $Row_{i+1} + Row_{M-1}$, etc.). After that, the processing continues until the last Row is processed.

The overall schematic is shown in Fig. 12.




Fig. 12. General schematic of the Kadane2D hardware

In this schematic, the control signals have been simplified to avoid overcharging and improve the readability of the design.

All components have been described in parameterizable VHDL code. Thus, the design becomes portable on any hardware support system. At the top level, the only generics that need to be specified by the user are M – the number of rows, and N – the number of columns of the input array.

6 Results of the 2D Implementation and Conclusions

The main goal of this paper was to create a hardware implementation of a well-known software algorithm, i.e. Kadane's algorithm, in both its 1D and 2D versions, for determining the maximal subsequence of a stream or an array of integers.

The paper presents a detailed step-by-step methodology for these adaptations. All choices made for the nature and the size of the used resources are discussed and justified.

Like for the 1D problem, the design was specified in parameterizable VHDL code and the target chips chosen as physical support of the implementation were Xilinx Virtex devices. The design was simulated and tested for three different families: Virtex II (XC2V500-5-FG256), Virtex II PRO (XC2VP7-5-FF896) and Spartan3 (XC3S500E-5-FG320). The reports generated by Xilinx synthesis tools are presented in Table 3, which summarizes their most important aspects.

Table 3. Results of the 2D hardware implementation

Family	Virtex2	Virtex2PRO	Spartan3
Device	xc2v500-5-fg256	xc2vp7-5-ff896	xc3s500e-5-fg320
Device utilization summary			
No. of Slices	116 - 3%	116 - 2%	120 - 2%
No. of Slice Flip Flops	135 - 2%	135 - 1%	135 - 1%
No. of 4 input LUTs	194 - 3%	194 - 1%	191 - 2%
No. of Bonded IOBs	15 - 8%	15 - 3%	15 - 6%
No. of BRAMs	2 - 6%	2 - 4%	2 - 10%
No. of GCLKs	1 - 6%	1 - 6%	1 - 4%
Timing summary			
Speed grade	-5	-5	-5
Minimum period (ns)	10.904	10.431	11.693
Maximum Frequency (MHz)	91.709	95.869	85.518
Minimum input arrival time before clock (ns)	3.182	3.133	4.240
Maximum output required time after clock (ns)	4.840	3.997	4.105

In fact, as can be seen in Fig. 12, the implementation for the 2D problem is an extension of the 1D design, based on the same working principle. The speed is basically identical, except for the small delays introduced by the need to HOLD and RESET components after extracting each Row from the Memory, as explained in Section 5. The working frequency is now lower than in the 1D implementation: around 90 MHz, varying with the FPGA family, as shown in Table 3.

The simulation was done in ModelSIM and the proper functioning of the design was confirmed.

In conclusion, the overall performance is *at least* two orders of magnitude better in hardware than in software. The space available in FPGA Virtex devices allows implementations for large input arrays, the main limitation being introduced by the amount of BlockRAM memory available in the Virtex devices. If a larger size memory is necessary, it can be implemented outside the chip, but in this case the working frequency will obviously be smaller.

Future work will involve creating a hardware algorithm for determining not only the maximal (or minimal) subsequence, but the first k maximal (or minimal) subsequences in 1D or 2D input arrays.

7 Acknowledgements

This work was supported by the Romanian Ministry of Education and Research, under grant AT 178 / 2006.

References:

- [1] Bentley, J., "Programming pearls: algorithm design techniques." *Commun. ACM* 27,9 (1984), p. 865.
- [2] Bentley, J., "Programming pearls: perspective on performance." *Commun. ACM* 27,11 (1984), p. 1087.
- [3] Sung Eun Bae and Tadao Takaoka, "Algorithms for the Problem of K Maximum Sums and a VLSI Algorithm for the K Maximum Sub-arrays Problem." *Proceedings of the 7th International Symposium on Parallel Architectures, Algorithms and Networks (ISPAN'04)*, 1087-4089/04, 2004.
- [4] Hamid Reza Naji and B. Earl Wells, "On Incorporating Multi Agents in Combined Hardware/Software based Reconfigurable Systems -- A General Architectural Framework". *Proceedings of the 7th International Symposium on Parallel Architectures, Algorithms and Networks (ISPAN'02)*, 2002.
- [5] M. Wooldridge, "Agent-based Software Engineering," *In IEE Proceedings on Software Engineering*, 144(1), pages 26--37, February 1997.
- [6] X. Meng, V. Chaudhary, Bio-sequence analysis with cradle's 3SoC™ software scalable system on chip, *Proceedings of the 2004 ACM symposium on Applied computing*.
- [7] K. Sullivan, J. Yang, D. Coppit, S. Khurshid, D. Jackson - *Software assurance by bounded exhaustive testing*, Proceedings of the 2004 ACM SIGSOFT international symposium on Software testing and analysis.